

# Scheduling

# Scheduling

- When more than one process is ready to run, but **only one CPU** is available, a choice is to make
- **Part** of **OS** that does it is **scheduler**
- The algorithm it uses is **scheduling algorithm**

# Importance of Scheduling

- **Good** scheduling algorithms can make a **big** difference
  - Resource utilization
  - Perceived performance & User satisfaction
  - Meeting other system goals (e.g., important tasks being taken care of immediately)

# When to Schedule

- When a new process is created:
  - Parent or child? Both are Ready
  - which one to run?
- When a process exits:
  - One of the ready processes should be run
- When a process blocks: Another process has to be selected to run
  - Blocking may occur for:
    - I/O
    - Semaphore

# When to Schedule



- When an I/O interrupt occurs:
  - In case of an interrupt of an I/O device having **completed** its work, some blocked process may now be ready
- If a h/w clock provides **periodic** interrupt:  
A scheduling decision can be made at each (or kth ) clock interrupt

# Preemptive & Non-preemptive

Classification of **Scheduling Algorithm** depending on dealing with clock interrupt

- **Non-preemptive**: Picks a process to run and lets it run until it **blocks** or voluntarily releases the CPU. **In effect at each clock interrupt, no scheduling is done.**
- **Preemptive**: Picks a process and lets it run for a maximum of some fixed time. If still running, it is **suspended** and another is picked.
- Preemptive scheduling requires having a **clock interrupt** occur at the end of the time interval to give **control** of the CPU back to the **scheduler**

# Batch Systems

- Common performance metrics
  - **Throughput**: number of jobs **completed** per hour 
  - **Turnaround time**: average time between the **submission** and **completion** of a job 
- Maximizing Throughput may not necessarily minimize Turnaround time

# First Come First Serve (FCFS)

- Process that requests the CPU FIRST is allocated the CPU FIRST.
- Also called FIFO
- **non**-preemptive
- Used in Batch Systems
- Real life analogy?
  - Transaction at Sonali Bank
- Implementation
  - FIFO queues
  - A **new** process enters the **tail** of the queue
  - The **schedule** selects from the **head** of the queue.



# FCFS Example

Process	Duration	Order	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

The final schedule:



P1 turnaround: 24

P2 turnaround: 27

P3 turnaround: 31


The average turnaround:

$$(24+27+31)/3 = 27.33$$

# FCFS Example 2

Process	Duration	Order	Arrival Time
P1	24	3	0
P2	3	1	0
P3	4	2	0

## The final schedule:



P1 turnaround: 31  
P2 turnaround: 3  
P3 turnaround: 7

The average turnaround:  
 $(31+3+7)/3 = 13.67$

# Advantage

- Easy to understand and implement
- Fair for equivalent processes

# Problems with FCFS

- Non-preemptive
- Non optimal turnaround
- Cannot utilize resources in parallel:
  - Assume 1 process CPU bounded and many I/O bounded processes
  - result: Convoy effect,
    - low CPU and I/O Device utilization
  - Why?

# Shortest Job First (SJF)

- Scheduling algorithm in **batch** systems
- Schedule the job with the shortest run time first
- Requirement: **the run time needs to be known in advance**
- SJF is **optimal** in terms of turnaround, if **all** jobs arrive at **same time**

# SJF: Example

Process	Duration	Order	Arrival Time
P1	6	2	0
P2	8	4	0
P3	7	3	0
P4	3	1	0




Do it yourself

P4 turnaround: 3  
P1 turnaround: 9  
P3 turnaround: 16  
P2 turnaround: 24

Total execution time: 24  
The average turnaround:  
 $(3+9+16+24)/4 = 13$

# Comparing to FCFS

Process	Duration	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0




P1 turnaround: 6  
P2 turnaround: 14  
P3 turnaround: 21  
P4 turnaround: 24

The total time is the same.  
The average turnaround:  
 $(6+14+21+24)/4 = 16.25$   
(comparing to 13)

# SJF is not always optimal

- SJF optimal only if all jobs have arrived at scheduling time

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 10

P2 turnaround: 10

The average turnaround (AWT):  
 $(10+10)/2 = 10$




# Preemptive SJF

- Also called **Shortest Remaining Time Next**
  - Schedule the job with the shortest **remaining** time required to complete
  - When **new** job **arrives**, compare its **total** time with the **remaining** time of the running job
  - If the new job needs less time the current job is suspended and the new job started
- Requirement: the run time needs to be known in advance

# Preemptive SJF: Same Example

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



P1 turnaround: 12  
P2 turnaround: 2

The average turnaround:  
 $(2+12)/2 = 7$

# Problem with Preemptive SJF?

- Starvation
  - In some condition, a job is waiting for ever
  - Example: Preemptive SJF
    - Process A with run time of 1 hour arrives at time 0
    - But every 1 minute, a short process with run time of 1 minute arrives
    - Result of Preemptive SJF: A never gets to run

# Interactive System



- Example: Servers
  - Serve multiple remote users all of whom are in a big hurry
- Performance Criteria
  - Min response time:
    - amount of time it takes from when a request was submitted until the **first response** is produced, not output
    - respond to requests quickly

# Interactive System

- Algorithms used here usually preemptive
  - Time is **sliced** into quantum (time intervals)
  - Scheduling decision is also made at the beginning of each quantum
- Representative algorithms:
  - Round-robin
  - Priority-based
  - Shortest process time
  - Guaranteed Scheduling
  - Lottery Scheduling
  - Fair Sharing Scheduling

# Round Robin

- Round Robin (RR)
  - Often used for timesharing
  - Each process is given a time slice called a *quantum*
  - It is run for the quantum or until it blocks
  - RR allocates the CPU uniformly (fairly) across participants from ready queue.
- Problem:
  - Do not consider priority
  - Context switch overhead




# Implementing Round Robin

- Keep the ready queue as a FIFO queue of processes.
- **New** processes are added to the **tail** of the ready queue.
- The scheduler
  - picks the **first** process from the ready queue
  - sets a timer to interrupt after 1 time quantum, and
  - Starts the process.
- When the quantum is over
  - The running process will be put at the **tail** of the ready queue.

# RR with Time Quantum = 20

<u>Process</u>	<u>Run Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- All processes arrive at time 0
- The **Gantt** chart is



- Higher average turnaround than SJF
- But better response time



# RR: Choice of Time Quantum

- Performance depends on length of the timeslice
  - Context switching isn't a free operation.
  - If timeslice time is set too high
    - attempting to amortize context switch cost, you get FCFS.
      - i.e. processes will finish or block before their slice is up anyway
      - Poor response time
  - If it's set too low
    - you're spending all of your time context switching between threads.

# Priority Scheduling

- Each job is assigned a priority
- Select **highest** priority job to run next
- Rational: higher priority jobs are more important
  - Example: simulation vs. auto save a document
- Problems:
  - Low priority process may **starve**
- Solution:
  - Priority need to be **adjusted** depending on the situation

# Assign Priority

- Two approaches
  - Static (for system with well known and regular application behaviors)
  - Dynamic (otherwise)
- Priority may be based on:
  - Cost to user.
  - Importance of user
  - Percentage of CPU time used in last X hours

# Example: **Dynamic** Priority Assignment

- Whenever highly I/O bound processes want the CPU it should be given the CPU immediately.
- Why?
- A simple algorithm for giving priority to I/O bound processes is to set the priority to  $1/f$ 
  - $f$  is the fraction of the last quantum used by a process
  - A process that used only 1 msec of its 50 msec quantum would get priority 50
  - A process that used 25 msec of its 50 msec quantum would get priority 2