

Interprocess Communication

Interprocess Communication

- Consider shell pipeline
 - `cat chapter1 chapter2 chapter3 | grep tree`
 - 2 processes
 - Information sharing
 - Order of execution


Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes require a **mechanism** to exchange data and information

IPC issues

1. How one process **passes** information to another ?
2. How to make sure that two or more processes do not get into each other's way when engaging in **critical** activities?
3. How to do proper **sequencing** when **dependencies** are present?
 - Ans 1: easy for threads, for processes different approaches (e.g., message passing, shared memory)
 - Ans 2 and Ans 3: same problems and same solutions apply for threads and processes
 - **Mutual exclusion & Synchronization**

Spooling Example: Correct




Spooling Example: Races

Process 1
int next_free;

- ① next_free = in;
- ③ Stores F1 into next_free;
- ④ in = next_free + 1

Shared memory



Process 2
int next_free;

- ② next_free = in
/* value: 7 */
- ⑤ Stores F2 into next_free;
- ⑥ in = next_free + 1

Better Coding?

- In previous code

```
for(;;){  
    int next_free = in;  
    slot[next_free] = file;  
    in = next_free+1;  
}
```

- What if we use one line of code?

```
for(;;){  
    slot[in++] = file  
}
```

When Can process Be switched?

- After each **machine** instruction!
- `in++` is a C/C++ statement, translated into **three** machine instructions:
 - load mem, R
 - inc R
 - store R, mem
- Interrupt (and hence process swiching) can happen in between.

Race condition

- Two or more processes are reading or writing some **shared** data and the final result depends on who runs precisely when
- Very hard to Debug

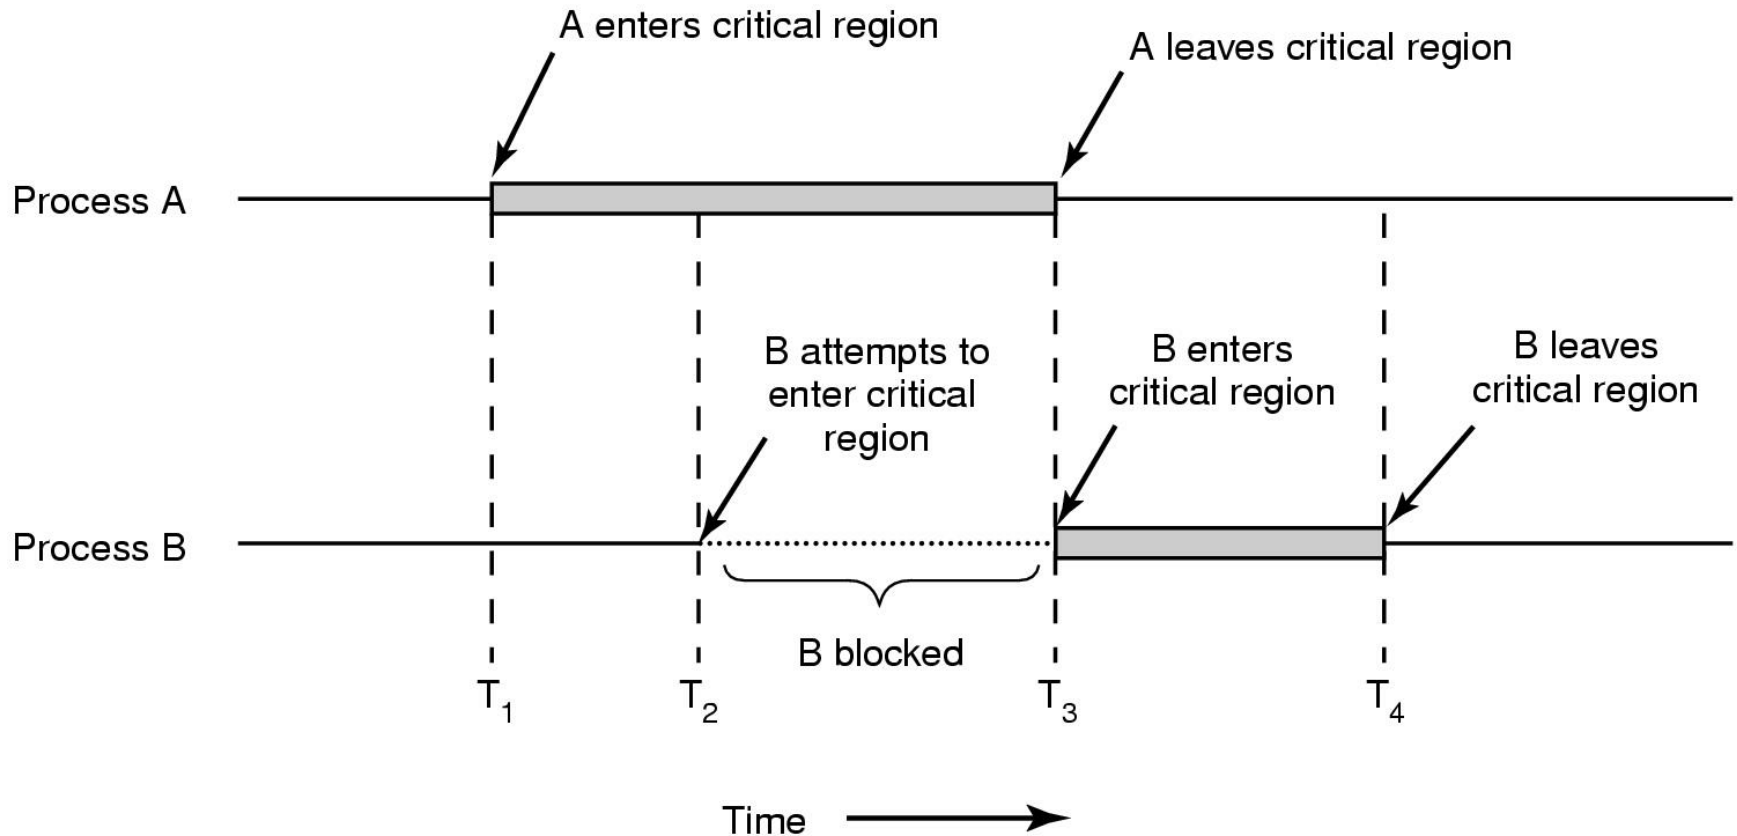
Critical Region

- That **part** of the program that do critical things such as accessing shared memory
- Can lead to race condition

Solution Requirement

- 1) No two processes **simultaneously** in **critical region**
- 2) No assumptions made about speeds or numbers of CPUs
- 3) No process running **outside** its critical region may **block** another process
- 4) No process must wait forever to enter its critical region

Solution Requirement



Sleep & wakeup

- When a process has to **wait**, change its state to **BLOCKED**
- Switched to **READY** state, when it is OK to retry entering the critical section
- Sleep is a **system call** that causes the caller to block
 - be suspended until another process wakes it up
- Wakeup system call has one parameter, the process to be awakened.
- Let's illustrate the use of sleep & wakeup with an example: **The producer consumer problem**

Producer Consumer Problem

- Also called bounded-buffer problem
- Two (or $m+n$) processes share a **common** buffer
- One (or m) of them is (are) **producer**(s): put(s) information in the buffer
- One (or n) of them is (are) **consumer**(s): take(s) information out of the buffer
- Trouble and solution
 - Producer wants to put but buffer **full**- Go to **sleep** and **wake up** when consumer takes one or more
 - Consumer wants to take but buffer **empty**- go to sleep and wake up when producer puts one or more

Sleep and Wakeup

```
#define N 100
int count = 0;
```

```
/* number of slots in the buffer */
/* number of items in the buffer */
```

```
void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        if (count == N) sleep();
        insert_item(item);
        count = count + 1;
        if (count == 1) wakeup(consumer);
    }
}
```

```
void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();
        item = remove_item();
        count = count - 1;
        if (count == N - 1) wakeup(producer);
        consume_item(item);
    }
}
```

Producer-consumer problem

Sleep and Wakeup: Race condition

- **Busy waiting problem is resolved but the following **race condition** exists**
- Unconstrained access to *count*
 - CPU is given to P just after C has **read** count to be 0 but not yet gone to sleep.
 - P calls wakeup
 - Result is **lost** wake-up signal
 - Both will sleep forever



Semaphores

- A new variable type
- A kind of **generalized** lock
 - First defined by Dijkstra in late 60s
 - Main synchronization **primitive** used in original UNIX
- Semaphores are like integers, except
 - No negative values
 - Only operations allowed are *up* and *down* – can't read or write value, except to set it initially

Semaphores: Types

- **Counting semaphore.**
 - The value can range over an unrestricted domain
- **Binary semaphore**
 - The value can range only between 0 and 1.
 - On some systems, binary semaphores are known as **mutex** locks as they provide mutual exclusion

Semaphores: Operation

- Operation “down”:
 - if value > 0 ; value-- and then continue.
 - if value = 0; process is put to sleep without completing the down for the moment
 - Checking the value, changing it, and possibly going to sleep, is all done as an **atomic** action.
- Operation “up”:
 - increments the value of the semaphore addressed.
 - If one or more process were sleeping on that semaphore, one of them is chosen by the system (e.g. at **random**) and is allowed to complete its *down*
 - The operation of incrementing the semaphore and waking up one process is also **indivisible**
 - No process ever blocks doing an *up*.

Semaphores: Atomicity

- Operations must be **atomic**
 - Two *down*'s together can't decrement value below zero
 - Similarly, process going to sleep in *down* won't miss wakeup from *up* – even if they both happen at same time

Producer & consumer

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

```
void producer(void)
{
    int item;
```

```
    while (TRUE) {
        item = produce_item( );
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
```

```
}
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

```
void consumer(void)
{
    int item;
```

```
    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item( );
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
```

```
}
```

Semaphores in Producer Consumer Problem: Analysis

- 3 semaphores are used
 - *full* (initially 0) for counting occupied slots
 - *Empty* (initially N) for counting empty slots
 - *mutex* (initially 1) to make sure that Producer and Consumer do not access the buffer at the same time.
- Here 2 uses of semaphores
 - Mutual exclusion (mutex)
 - Synchronization (full and empty)
 - To guarantee that certain event sequences do or do not occur

Block on:	Unblock on:
Producer: insert in full buffer	Consumer: item inserted
Consumer: remove from empty buffer	Producer: item removed

Semaphores: Usage

1. Mutual exclusion
2. Controlling access to limited resource
3. Synchronization

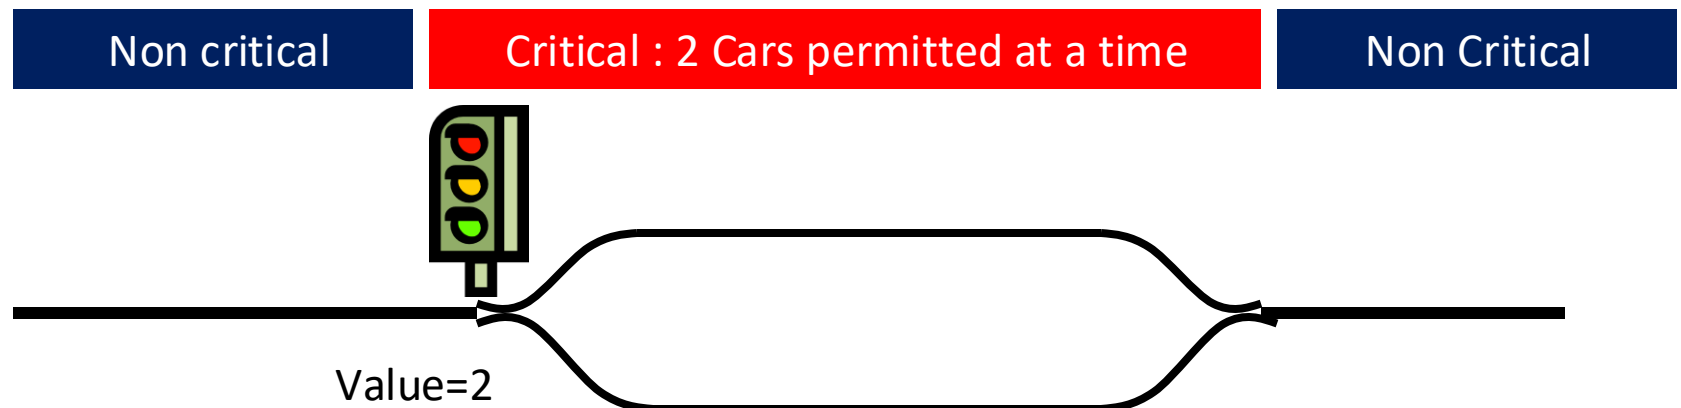
Mutual exclusion

- How to ensure that only one process can enter its C.R.?
- Binary semaphore **initialized to 1**
- Shared by all collaborating processes
- If each process does a *down* just before entering CR and an *up* just after leaving then mutual exclusion is guaranteed

```
do {  
    down (mutex) ;  
    // critical section  
    up (mutex) ;  
    // remainder section  
} while (TRUE);
```


Controlling access to a resource

- What if we want maximum **m** process/thread can use a resource simultaneously ?
- Counting semaphore **initialized to the number of available resources**
- Semaphore from railway analogy
 - Here is a semaphore **initialized to 2** for resource control:



Synchronization

- How to resolve dependency among processes
- Binary semaphore **initialized to 0**
- consider 2 concurrently running processes:
 - P1 with a statement S1 and
 - P2 with a statement S2.
 - Suppose we require that S2 be executed only after S1 has completed.

P1


```
S1;  
up(synch);
```

P2

```
down(synch);  
S2;
```

Semaphores: “Be Careful”

- Suppose the following is done in **Producer's** code



- If buffer **full** P would block due to `down(&empty)` with `mutex = 0`.
- So now if C tries to access the buffer, it would block too due to its `down(&mutex)`.
- Both processes would stay blocked **forever**:
DEADLOCK